

```

named `train` and `validation` with 80% and 20% of the data, respectively.
"""

data_len = data.shape[0]

train_indices = shuffled_indices[:pointeight]

return train, validation
train, validation = train_val_split(training_val_data)

```

Let's fit our linear regression model using the ordinary least squares estimator! We will start with something simple by using only two features: the **number of bedrooms** in the household and the **log-transformed total area covered by the building** (in square feet).

Consider the following expression for our first linear model that contains one of the features:

$$\text{Log Sale Price} = \theta_0 + \theta_1 \cdot (\text{Bedrooms})$$

In parallel, we will also consider a second model that contains both features:

$$\text{Log Sale Price} = \theta_0 + \theta_1 \cdot (\text{Bedrooms}) + \theta_2 \cdot (\text{Log Building Square Feet})$$

```

[13]: from feature_func import *      # Import functions from Project A1

##### Copy any function you would like to below #####
...
#####

def feature_engine_simple(data):
    # Remove outliers

    # Create Log Sale Price column

    # Create Bedroom column

    # Select X and Y from the full data
    X = data[['SqrFt', 'Bedroom']]
    Y = data[['Log Sale Price']]
    return X, Y

# Reload the data
full_data = pd.read_csv("cook_county_train.csv")

# Process the data using the pipeline for the first model.

```

```
np.random.seed(1337)
train_m1, valid_m1 = train_test_split(X, Y, test_size=0.2,
                                       random_state=1337)
X_train_m1_simple, X_valid_m1_simple, Y_train_m1_simple, Y_valid_m1_simple =
```

```
# Take a look at the result
display(X_train_m1_simple.head())
display(Y_train_m1_simple.head())
```

	Bedrooms
130829	4
193890	2
30507	2
91308	2
131132	3

130829	12.994530
193890	11.848683
30507	11.813030
91308	13.060488
131132	12.516861

Name: Log Sale Price, dtype: float64

```

# Helper function
def select_columns(data, *columns):
    """Select only columns passed as arguments."""

# Pipelines, a list of tuples
m1_pipelines = [
    (remove_outliers, None, {
    },
    (log_transform,
    (add_total_bedrooms,
    (select_columns, ['L
]

X_train_m1, Y_train_m1

X_valid_m1, Y_valid_m1 =

# Take a look at the result
# It should be the same above as the result returned by feature_engine_simple
display(X_train_m1.head())
display(Y_train_m1.head())

```

	Bedrooms
130829	4
193890	2
30507	2
91308	2
131132	3
130829	12.994530
193890	11.848683
30507	11.813030
91308	13.060488
131132	12.516861

Name: Log Sale Price, dtype: float64

```

m2_pipelines =
    [
        Pipeline([
            FeatureSelector(
                lower=100,
            ),
            (log_transform,
            )
        ])
    ]

```

```

X_train_m2, Y_train_m2 =
    train_test_split(X_train, Y_train,
                    test_size=0.2,
                    random_state=42,
                    stratify=Y_train)
X_valid_m2, Y_valid_m2 =
    train_test_split(X_valid, Y_valid,
                    test_size=0.2,
                    random_state=42,
                    stratify=Y_valid)

```

```

# Take a look at the result
display(X_train_m2.head())
display(Y_train_m2.head())

```

	Bedrooms	Log Building Square Feet
130829	4	7.870166
193890	2	7.002156
30507	2	6.851185
91308	2	7.228388
131132	3	7.990915
130829	12.994530	
193890	11.848683	
30507	11.813030	
91308	13.060488	
131132	12.516861	

Name: Log Sale Price, dtype: float64


```
[17]: grader.check("q3b")
```


1.14 Question 4a

One way of understanding a model's performance (and appropriateness) is through a plot of the residuals versus the observations.

```
[22]: plt.scatter
      plt.xlabel(
      plt.ylabel(
      plt.title('Residuals between Validation Log Sale Price v Predicted Log Sale_
      ↪Price') 8)
```

```
[22]: Text(0.5, 1.0, 'Residuals between Validation Log Sale Price v Predicted Log Sale
Price')
```





```
[27]: full_data.describe()[["Estimate (Land)"]] ## high should be around 200,000k
```

```
[27]:      Estimate (Land)
count    2.047920e+05
mean      5.187066e+04
std       5.591360e+04
min        0.000000e+00
25%       2.762000e+04
50%       3.795000e+04
75%       5.580000e+04
max       3.716680e+06
```

```
[28]: full_data["Estimate (Land)"].describe(percentiles=[0.95])
```

```
[28]: count    2.047920e+05
mean      5.187066e+04
std       5.591360e+04
min        0.000000e+00
50%       3.795000e+04
95%      1.320445e+05
max       3.716680e+06
```


Name: Estimate (Land), dtype: float64

```
[29]: full_data["Sale Price"].describe(percentiles=[0.95]) # high should be 700k
```

```
[29]: count    2.047920e+05  
      mean    2.451646e+05  
      std     3.628694e+05  
      min     1.000000e+00  
      50%     1.750000e+05  
      95%     7.750000e+05  
      max     7.100000e+07  
      Name: Sale Price, dtype: float64
```

```
[30]: full_data_cop = full_data.copy()  
      full_data_cop["Log Sale Price"] = np.log(full_data_cop["Sale Price"])  
      full_data_cop["ages"] = np.log(full_data_cop["Age"])  
      plt.scatter(data = full_data_cop, x = "ages", y = "Log Sale Price")
```

```
[30]: <matplotlib.collections.PathCollection at 0x7fd101641050>
```



```
[31]: np.percentile
```


`LinearRegression()` model with intercept term for grading purposes. Do not modify any hyperparameter in `LinearRegression()`, and focus on feature selection or hyperparameters of your own feature engineering function.

It may also be helpful to calculate the RMSE directly as follows:

$$RMSE = \sqrt{\frac{\sum_{\text{houses in the set}} (\text{actual price for house} - \text{predicted price for house})^2}{\text{number of houses}}}$$

A function that computes the RMSE is provided below. Feel free to use it if you would like calculate the RMSE for your training set.

```
[35]: def rmse(predicted, actual):  
    """  
    Calculates RMSE from actual and predicted values.  
    Input:  
        predicted (1D array): Vector of predicted/fitted values  
        actual (1D array): Vector of actual values  
    Output:  
        A float, the RMSE value.  
    """  
    return np.sqrt(np.mean((actual - predicted)**2))
```

```

[41]: True Log Sale Price Predicted Log Sale Price True Sale Price \
1      12.560244      11.966303      285000.0
2      9.998798      11.475456      22000.0
3     12.323856     11.790922     225000.0
4     10.025705     11.354999      22600.0
6     11.512925     12.153325     100000.0

Predicted Sale Price
1     157361.864168
2      96322.366007
3     132048.214772
4      85391.220473
6     189723.911211

```

The lower interval contains houses with true sale price \$735.0 to \$219696.0
The higher interval contains houses with true sale price \$219696.0 to \$1202604.0

```
[43]: rmse_cheap =  
      rmse_expensiv
```

```
prop_overest_cheap =  
    1 - 105.7  
prop_overest_expensive =
```

```
print(f"The RMSE for properties with log sale prices in the interval_  
↳ {(min_Y_true, median_Y_true)} is {np.round(rmse_cheap)}")  
print(f"The RMSE for properties with log sale prices in the interval_  
↳ {(median_Y_true, max_Y_true)} is {np.round(rmse_expensive)}\n")  
print(f"The percentage of overestimated values for properties with log sale_  
↳ prices in the interval {(min_Y_true, median_Y_true)} is {np.round(100 *_  
↳ prop_overest_cheap, 2)}%")  
print(f"The percentage of overestimated values for properties with log sale_  
↳ prices in the interval {(median_Y_true, max_Y_true)} is {np.round(100 *_  
↳ prop_overest_expensive, 2)}%")
```

The RMSE for properties with log sale prices in the interval is

The RMSE for properties with log sale prices in the interval) is

The percentage of overestimated values for properties with log sale prices in the interval is

The percentage of overestimated values for properties with log sale prices in the interval is

